

# Pointer Structures

Troels Henriksen ([athas@sigkill.dk](mailto:athas@sigkill.dk))

DIKU  
University of Copenhagen

Arrays and Records

Linked Lists


Trees

## Representing arrays of tuples

Consider arrays of type `[](i32, i8)`. Since an `i32` is four bytes and a `i8` is one byte, how is this stored in memory?

## Representing arrays of tuples


Consider arrays of type `[](i32, i8)`. Since an `i32` is four bytes and a `i8` is one byte, how is this stored in memory?



**Problem?**

## Representing arrays of tuples

Consider arrays of type `[](i32, i8)`. Since an `i32` is four bytes and a `i8` is one byte, how is this stored in memory?




**Problem?** Unaligned accesses.

## Representing arrays of tuples

Consider arrays of type `[](i32, i8)`. Since an `i32` is four bytes and a `i8` is one byte, how is this stored in memory?




**Problem?** Unaligned accesses.




**Problem?**

## Representing arrays of tuples

Consider arrays of type `[](i32, i8)`. Since an `i32` is four bytes and a `i8` is one byte, how is this stored in memory?



**Problem?** Unaligned accesses.



**Problem?** Waste of memory.

# Tuples of arrays

## Representation

An array  $[](t_1, t_2, t_3 \dots)$  is represented in memory as  $([]t_1, []t_2, []t_3 \dots)$ , i.e. as *multiple arrays*, each containing only primitive values.

0	1	2	3	4	5	6	7	8	9	10
i32				i32				i32	...	
i8	i8	i8	i8	i8	i8	i8	i8	i8	i8	...

- Common (and crucial) optimisation.
- Called “struct of arrays” in legacy languages.
- Automatically done by the Futhark compiler.

## "Unzipped" SOACs

Instead of

```
let tmp = map (\(x,y) -> (x-1, y+1))  
    (zip xs ys)  
let (xs, ys) = unzip xs_ys'
```

could we write

```
let (xs, ys) = map (\x y -> (x-1, y+1)) xs ys
```

?


- Annoying to give type rules, **but this is actually what the compiler does internally.**
- **Isomorphic to source language**, but this form is easier to manipulate in a compiler.
- This will be relevant in a later lecture.

Arrays and Records

Linked Lists

Trees

## Standard representation: Cells with Pointers




```
data List a = Nil  
            | Cons a (List a)
```

-- values look like

```
Cons 1 (Cons 2 (Cons 3 Nil))
```

## Standard representation: Cells with Pointers



```
data List a = Nil  
            | Cons a (List a)
```

-- values look like  
Cons 1 (Cons 2 (Cons 3 Nil))

### Challenges for Data Parallelism

- Many languages do not support recursive data structures.
- Traversing a list is **sequential**.

Let us try to address these problems.

## An array encoding of the link structure for list with $n$ nodes

The *Successor Array*  $S$  of length  $n$

$S[i]$  denotes index of successor for node  $i$ , with  $S[i] = n$  indicating last element.

The *Value Array*  $V$  of length  $n$

$V[i]$  denotes the value of node  $i$ .

Example

$$S = [4, 0, 5, 2, 3]$$

encodes a list with nodes stored in order

$$[1, 0, 4, 3, 2]$$

meaning the first node is at index 1, second node at index 0, etc.

# Questions

$$S = [4, 0, 5, 2, 3]$$

- Given  $S$ , how do we find the head node?

# Questions

$$S = [4, 0, 5, 2, 3]$$

- Given  $S$ , how do we find the head node?
  - ▶ Must find the node *not* marked as successor by any other node.
  - ▶ Can be done with scatter followed by reduce.

# Questions

$$S = [4, 0, 5, 2, 3]$$

- Given  $S$ , how do we find the head node?
  - ▶ Must find the node *not* marked as successor by any other node.
  - ▶ Can be done with scatter followed by reduce.
- How do we find the last node?

# Questions

$$S = [4, 0, 5, 2, 3]$$

- Given  $S$ , how do we find the head node?
  - ▶ Must find the node *not* marked as successor by any other node.
  - ▶ Can be done with scatter followed by reduce.
- How do we find the last node?
  - ▶ Find the node  $i$  where  $S[i] = n$ .
  - ▶ Just a reduce.

# Questions

$$S = [4, 0, 5, 2, 3]$$

- Given  $S$ , how do we find the head node?
  - ▶ Must find the node *not* marked as successor by any other node.
  - ▶ Can be done with scatter followed by reduce.
- How do we find the last node?
  - ▶ Find the node  $i$  where  $S[i] = n$ .
  - ▶ Just a reduce.

**What about all the other nodes?**

# List Ranking

## The List Ranking Problem

Determine the rank of each element in list, such that first element has rank 1, second has rank 2, etc.

# List Ranking

## The List Ranking Problem

Determine the rank of each element in list, such that first element has rank 1, second has rank 2, etc.

- We will actually study a variant, where the *last* element has rank 1.
- Can see it as “distance from end”.
- Are we cheating? **Does that make the problem harder or easier?**

## Wyllie's List Ranking

Each list element  $i$  has a *successor pointer*  $S[v]$ , and at each time step we update

$$S[v] \leftarrow \begin{cases} S[S[v]] & \text{when } S[v] \neq n \\ S[v] & \end{cases}$$

- Distance covered by pointer doubles for each time step.
- After  $\lceil \log(n) \rceil$  steps each  $S[i]$  is to  $n$  (end of list).
- Total work is  $O(n \log(n))$  and span  $O(\log(n))$ .

## Wyllie's List Ranking

Each list element  $i$  has a *successor pointer*  $S[v]$ , and at each time step we update

$$S[v] \leftarrow \begin{cases} S[S[v]] & \text{when } S[v] \neq n \\ S[v] & \end{cases}$$

- Distance covered by pointer doubles for each time step.
- After  $\lceil \log(n) \rceil$  steps each  $S[i]$  is to  $n$  (end of list).
- Total work is  $O(n \log(n))$  and span  $O(\log(n))$ .

Can compute rank by initially setting  $R[i] = 1$  and then

$$R[v] \leftarrow \begin{cases} R[i] = R[i] + R[S[i]] & \text{when } S[v] \neq n \\ R[i] & \end{cases}$$

in each step.

## In Futhark

```
def step [n] (R: [n]i32) (S: [n]i64) =
  let f i = if S[i] == n
            then (R[i], S[i])
            else (R[i] + R[S[i]]), S[S[i]])
  in unzip (tabulate n f)

def wyllie [n] (S: [n]i64) : [n]i32 =
  let R = replicate n 1
  let (R,_) = loop (R, S) for _i < 64 - i64.clz n do
    step R S
  in R
```

Is this work efficient?

## In Futhark

```
def step [n] (R: [n]i32) (S: [n]i64) =
  let f i = if S[i] == n
            then (R[i], S[i])
            else (R[i] + R[S[i]]), S[S[i]])
  in unzip (tabulate n f)

def wyllie [n] (S: [n]i64) : [n]i32 =
  let R = replicate n 1
  let (R,_) = loop (R, S) for _i < 64 - i64.clz n do
    step R S
  in R
```

Is this work efficient?

- No, a sequential implementation has work  $O(n)$ .
- Reason is that we keep inspecting nodes that have already finished.
- Work efficient algorithms exist, but are more complicated.

## Converting list to array

```
def list_to_array [n] 'a (V: [n]a) (S: [n]i64) =  
  scatter (copy V) (map (\i -> n - i64.i32 i) (wyllie S)) V
```

## Converting list to array

```
def list_to_array [n] 'a (V: [n]a) (S: [n]i64) =  
  scatter (copy V) (map (\i -> n - i64.i32 i) (wyllie S)) V
```

- To scan or reduce a list we could convert to array, then use array operations.
- **Is there another option?**

## List ranking, but now a scan

```
def wyllie_scan_step [n] 'a (op: a -> a -> a)
                        (V: [n]a) (S: [n]i64) =
let f i = if S[i] == n
            then (V[i], S[i])
            else (V[i] `op` V[S[i]], S[S[i]])
in unzip (tabulate n f)

def wyllie_scan [n] 'a (op: a -> a -> a)
                        (V: [n]a) (S: [n]i64) =
let (V,_) = loop (V, S) for _i < 64 - i64.clz n do
    wyllie_scan_step op V S
in V
```

## Packing it all up

What might it look like to actually construct a library of list operations?

```
type list [n] 'a = { S: [n]i64
                      , V: [n]a
                      , head: i64
                      , last: i64
                    }

def head [n] 'a (l: list [n] a) : a = l.V[l.head]

def last [n] 'a (l: list [n] a) : a = l.V[l.last]
```

## From array to list

## From array to list

```
def from_array 'a [n] (V: [n]a) : list [n] a =
{ S = map (+1) (iota n)
, V
, head = 0
, last = n-1
}
```

## Reversing a list

## Reversing a list

```
def rev [n] 'a (l: list [n] a) =  
  let f i = (l.S[i], i)  
  let (is, vs) = unzip (tabulate n f)  
  in l with S = scatter (replicate n n) is vs  
    with head = l.last  
    with last = l.head
```

# Scans

```
def scan [n] 'a (op: a -> a -> a) (l: list [n] a) =
  let l' = rev l
  in l with V = (wyllie_scan op l'.V l'.S)
```

## What about reductions?

```
def reduce [n] 'a (op: a -> a -> a)
           (_ne: a) (l: list [n] a) =
    last (scan op l)
```

**Note:** *list* last, not the builtin array last.

## What about reductions?

```
def reduce [n] 'a (op: a -> a -> a)
           (_ne: a) (l: list [n] a) =
    last (scan op l)
```

**Note:** *list* last, not the builtin array last.

We can do better in the *commutative* case.

```
def reduce_comm [n] 'a (op: a -> a -> a)
                   (ne: a) (l: list [n] a) =
reduce_comm op ne l.V -- Array reduction.
```

# Concatenation

# Concatenation

```
def (++) [n] [m] 'a (x: list [n] a) (y: list [m] a) =
{ S = if n == 0 || m == 0 then x.S ++ y.S
  else (copy x.S with [x.last] = n + y.head)
        ++
        map (\i -> i+n) y.S
  , V = x.V ++ y.V
  , head = x.head
  , last = y.last + n
}
```

# Take?

```
val take [n] 'a (i: i64) (l: list [n] a) : list [i] a
```

*Exercise for the reader.*

Arrays and Records


Linked Lists

Trees


## Trees are also pointer structures

- Optimal representation depends on what we want to do with them.
- The “parent pointer representation” will be our main object of study, but we will look at a few others, too.
- As with lists, the main problem is to find a **linearized representation**.

# Semi-linear representations




- Could represent this as a linked list (or irregular array) of *levels*.



- How do we know the parent-child relationship?
  - ▶ Can add ancillary structure.
  - ▶ But this is still *recursive*, and we can't have that.


## A representation for binary trees

- Child vectors  $L, R$  where  $L[i]$  is the left child of  $i$  and  $R[i]$  is the right child.
- $L[i] = -1$  or  $R[i] = -1$  denotes no such child.



## A representation for binary trees


- Child vectors  $L, R$  where  $L[i]$  is the left child of  $i$  and  $R[i]$  is the right child.
- $L[i] = -1$  or  $R[i] = -1$  denotes no such child.



$$\begin{aligned}L &= [ \quad 1, \quad -1, \quad -1, \quad 4, \quad -1, \quad -1, \quad -1 \quad ] \\R &= [ \quad 3, \quad \quad 2, \quad -1, \quad 5, \quad -1, \quad -1, \quad -1 \quad ]\end{aligned}$$

This is an OK representation for many purposes, but it cannot represent  $n$ -ary trees.

## What if we flip the pointers?




### Parent Vector $P$

- $P[i]$  is the parent of  $i$ .
- Root is its own parent:  $P[i] = i$ .

Either siblings are considered unordered, or they are ordered by their index.

# What if we flip the pointers?



## Parent Vector $P$

- $P[i]$  is the parent of  $i$ .
- Root is its own parent:  $P[i] = i$ .

Either siblings are considered unordered, or they are ordered by their index.

## Converting from child vectors to parent vector

$$L[i] = j \vee R[i] = j \Rightarrow P[j] = i$$

## A more complicated tree

Can you draw this tree?

$$P = [0, 0, 0, 0, 0, 3, 3, 4]$$


## A more complicated tree

Can you draw this tree?

$$P = [0, 0, 0, 0, 0, 3, 3, 4]$$

What about this one?


$$P = [1, 7, 7, 2, 2, 7, 7, 7]$$



# A more complicated tree


Can you draw this tree?

$$P = [0, 0, 0, 0, 0, 3, 3, 4]$$



What about this one?


$$P = [1, 7, 7, 2, 2, 7, 7, 7]$$



## A more complicated tree


Can you draw this tree?

$$P = [0, 0, 0, 0, 0, 3, 3, 4]$$



What about this one?


$$P = [1, 7, 7, 2, 2, 7, 7, 7]$$



- The same tree can have different parent vectors.
- **Element order (almost) does not matter!**

## Depth Vectors

For a tree  $P = [0, 0, 0, 0, 0, 3, 3, 4]$



we can state the distance of each node from the root as

$$D = [0, 1, 1, 1, 1, 2, 2, 2]$$

- Monotonically increasing here, but that depends on node ordering.


### Depth Vector $D$

$D[i]$  is the distance of node  $i$  from the root.

## Using depth vector as representation

By assuming a specific node ordering (e.g. preorder traversal) the depth vector is an unambiguous representation of the tree.

For a tree




the preorder traversal is

$$[0, 1, 2, 3, 5, 6, 4, 7]$$

and so

$$D = [0, 1, 1, 1, 2, 2, 1, 2]$$

## Constructing depth vector from parent vector




$$P = [0, 0, 0, 0, 0, 3, 3, 4]$$

Idea:

- The parent vector encodes *multiple linked lists* from leaves to the root.
  - ▶ [1, 0]
  - ▶ [2, 0]
  - ▶ [5, 3, 0]
  - ▶ [6, 3, 0]
  - ▶ [7, 4, 0]

## Constructing depth vector from parent vector




$$P = [0, 0, 0, 0, 0, 3, 3, 4]$$

### Idea:


- The parent vector encodes *multiple linked lists* from leaves to the root.
  - ▶ [1, 0]
  - ▶ [2, 0]
  - ▶ [5, 3, 0]
  - ▶ [6, 3, 0]
  - ▶ [7, 4, 0]
- Do essentially list ranking on each, simultaneously!
- Work  $O(n \log n)$ , span  $O(\log n)$ .

## Constructing parent vector from depth vector, assuming preorder



$$D = [0, 1, 1, 1, 2, 2, 1, 2]$$

## Constructing parent vector from depth vector, assuming preorder



$$D = [0, 1, 1, 1, 2, 2, 1, 2]$$


### Idea

- For each node, linear search *left* until you find a node with lower depth.
  - ▶ That is your parent.
- Potentially costly *sequential* search.
  - ▶ Span  $O(n)$  in worst case.
- **Optimisation idea.**
  - ▶ Pair depth vector with original index.
  - ▶ Sort in an appropriate way.
  - ▶ Use binary search.

# From traversal vector to depth vector

A traversal vector describes a preorder traversal of tree

- Contains elements 1 and  $-1$ .
  - $1$  descends into a new node.
  - $-1$  returns to parent.




[1,  $-1$ , 1,  $-1$ , 1, 1, 1,  $-1$ , 1,  $-1$ ,  $-1$ , 1, 1]

# From traversal vector to depth vector

A traversal vector describes a preorder traversal of tree

- Contains elements 1 and  $-1$ .
  - $1$  descends into a new node.
  - $-1$  returns to parent.



[1,  $-1$ , 1,  $-1$ , 1, 1,  $-1$ , 1,  $-1$ ,  $-1$ , 1, 1]


What happens if we take the prefix sum?

[1, 0, 1, 0, 1, 2, 1, 2, 1, 0, 1, 2]

# From traversal vector to depth vector

A traversal vector describes a preorder traversal of tree

- Contains elements 1 and  $-1$ .
  - $1$  descends into a new node.
  - $-1$  returns to parent.



[1,  $-1$ , 1,  $-1$ , 1, 1, 1,  $-1$ , 1,  $-1$ ,  $-1$ , 1, 1]

What happens if we take the prefix sum?

[1, 0, 1, 0, 1, 2, 1, 2, 1, 0, 1, 2]

*It is (almost) the depth vector!*

- Indexed from 1 instead of 0.
- Must remove the elements corresponding to  $-1$  in traversal.

## Summary of trees: the parent pointer representation

$n$ -node tree is represented by the these  $n$ -element vectors.

$P$ : parent vector.

- $P[i]$  is the parent of node  $i$ .
- For the root node,  $P[i] = i$ .

$D$ : depth vector.

- $D[i]$  is the distance from root to node  $i$ .
- For the root node,  $D[i] = 0$ .
- For other nodes,  $D[i] = D[P[i]] + 1$ .

$V$ : value vector.

- $V[i]$  is the value of node  $i$ .

- $P$  can be computed from  $D$ .
- $D$  can be computed from  $P$ .
- ...but usually convenient to have both.

## Summary

- Recursive pointer structures are not natural in data parallel languages, but can work well if we are careful.
- <https://github.com/diku-dk/containers/blob/main/lib/github.com/diku-dk/containers/list.fut>
- An interesting DPP project might be to implement various pointer structures (lists, trees, graphs) in a data parallel language and see what their performance is like.