

# Weekly Assignment 1

## Data Parallel Programming

Troels Henriksen  
Includes work by Martin Elsman  
DIKU, University of Copenhagen

November 2025

### Introduction

This weekly assignment aims at exercising the use of Futhark for writing parallel programs in a functional setting. The exercises assume access to a computer with Futhark installed. For information about installing Futhark, please consult <https://futhark-lang.org>.

The handin is expected to consist of a report in either plain text or PDF file (the latter is recommended unless you know how to perform sensible line wrapping by hand) of 4–6 pages, excluding any figures, along with an archive containing your source code. The report should contain instructions on how to run and benchmark your code.

# 1 Getting started

This exercise shows how to solve parallel problems in Futhark.

## Exercise 1.1: Write a function

Create a Futhark function called `process` that takes as arguments two one-dimensional `i32` arrays (signals) of the same length and computes the maximum absolute difference (pointwise) between the signals (you should not use Futhark's `loop` construct). The function should return the value 0 if two empty signals are passed to the function.

Consider the following two signals:

```
def s1 = [23,45,-23,44,23,54,23,12,34,54,7,2, 4,67]
def s2 = [-2, 3, 4,57,34, 2, 5,56,56, 3,3,5,77,89]
```

What is the result of calling your function on `s1` and `s2`?

**Hint** Use `futhark repl` to load the function into the Futhark interpreter where you can easily run it interactively.

## Exercise 1.2: Benchmark your function

Declare an entry point function `test_process` (with `entry` instead of `let`) that calls `process`. Then write a test block, as discussed in the Futhark book<sup>1</sup> and in class, that tests this function with random arrays of various sizes. The lengths of the arrays should be 100, 1000, 10000, 100000, 1000000, and 10000000.

Benchmark your program with `futhark bench` using both the `c`, `multicore`, and `openccl` backends. Map the timings (in microseconds) onto a chart and remember to specify the system on which you're running the executables.

## Exercise 1.3: Extend your function

Create a version of `process`, called `process_idx`, that also returns the index of the source signals for which the largest absolute difference is found.

Report the result of calling `process_idx` on the signals `s1` and `s2`. Show evidence that your solution scales as the `process` function. Do this by writing a `test_process_idx` function and a test block, as before.

**Hint** The lecture slides should give you a hint to solving this problem.

---

<sup>1</sup><https://futhark-book.readthedocs.io/en/latest/practical-matters.html>

## 2 Implementing prefix sum

For this question you will be implementing the Hillis-Steele prefix sum and the work-efficient prefix sum discussed in class. The implementations will be specialised to summation of `i32` values, and you may assume that the input is of size  $n = 2^m$ . Given such an  $n$ , you can obtain  $m$  by the following function.

```
def ilog2 (x: i64) = 63 - i64.i32 (i64.clz x)
```

Base your solution on the following skeleton:

```
def hillis_steele [n] (xs: [n]i32) : [n]i32 =
  let m = ilog2 n
  in loop xs = copy xs for d in ... do
    ...
    ...

def work_efficient [n] (xs: [n]i32) : [n]i32 =
  let m = ilog2 n

  let upswept =
    loop xs = copy xs for d in ... do
    ...
    let upswept[n-1] = 0

  let downsweped =
    loop xs = upswept for d in ... do
    ...
    in downsweped
```

Benchmark your two functions with `futhark bench --backend=opencl` (or `cuda` or `hip`) as shown in the lecture slides (and you did in the first task) and the Futhark book, and evaluate whether their actual performance matches the asymptotic guarantees. Compare their performance with the built-in `scan (+) 0i32`.

**Hint** Use Futhark's range notation to construct the iteration spaces for `d`. For example, `x..x-1...y` creates a decreasing sequence with stride 1 from `x` to `y` (both inclusive), and `x...y` creates an increasing sequence from `x` to `y` (both inclusive).

### 3 Segmented operations

#### Exercise 3.1: Proof of associativity

Assuming  $\oplus$  is an associative operator with neutral element 0, show that  $(0, \text{false})$  is a left-neutral element of

$$(v_1, f_1) \oplus' (v_2, f_2) = (\text{if } f_2 \text{ then } v_2 \text{ else } v_1 \oplus v_2, f_1 \vee f_2)$$

#### Exercise 3.2: Segmented scans and reductions

The operator in the previous question can be used to implement a segmented scan. Specifically, a `scan` on an array of type  $[]t$  with operator  $\oplus$  and neutral element  $0_{\oplus}$  can be turned into a segmented scan on an array of type  $[n](t, \text{bool})$  with operator  $\oplus'$  and neutral element  $(0_{\oplus}, \text{false})$ . A `true` indicates the beginning of a segment, and `false` the continuation of a segment.

Finish the following Futhark definition of segmented scan:

```
def segscan [n] 't (op: t -> t -> t) (ne: t)
    (arr: [n](t, bool)): [n]t =
    ...
```

A segmented reduction is more complicated, but can be implemented by first performing a segmented scan, and then making use of `scatter`.

Finish the following Futhark definition of segmented reduction:

```
def segreduce [n] 't (op: t -> t -> t) (ne: t)
    (arr: [n](t, bool)): []t =
    ...
```

Note that we cannot provide the size of the returned array in the type, as we do not know the number of segments.

Benchmark the performance of segmented scan versus ordinary scan, and segmented reduce versus ordinary reduction, and show the result.

### Exercise 3.3: Implementing `reduce_by_index`

Finish the following implementation of `reduce_by_index`:

```
def reduce_by_index 'a [m] [n]
    (dest : *[m] a)
    (f : a -> a -> a) (ne : a)
    (is : [n] i64) (as : [n] a) : *[m] a =
    ...

```

Obviously, do not use any of the built-in functions `hist` or `reduce_by_index` (nor its higher-dimensional variants). Instead, use `radix_sort_int_by_key`<sup>2</sup> from the `github.com/diku-dk/sorts` package and the `segreduce` you wrote previously.

**Hint** To use the `github.com/diku-dk/sorts` package, run

```
$ futhark pkg add github.com/diku-dk/sorts
$ futhark pkg sync
```

This will create a directory tree `lib/`, from which you can import the sorting function with

```
import "lib/github.com/diku-dk/sorts/radix_sort"
```

**Hint** After having sorted the values, you can use the `rotate` function together with `map2` to create a flag vector, which can be used as input to a call to `segreduce`.

What is the asymptotic complexity (work and span) of your implementation? Do you think this is optimal? How does it perform in practice compared to Futhark's built-in `reduce_by_index`?

---

<sup>2</sup>[https://futhark-lang.org/pkgs/github.com/diku-dk/sorts/0.4.2/doc/lib/github.com/diku-dk/sorts/radix\\_sort.html#5982](https://futhark-lang.org/pkgs/github.com/diku-dk/sorts/0.4.2/doc/lib/github.com/diku-dk/sorts/radix_sort.html#5982)

## 4 2D Ising Model

The 2D Ising Model is a mathematical modeling of the behaviour of a simple idealised ferromagnet, in which we compute the *spin* (roughly, polarity) of a grid of electrically charged atoms over a period of time. From a Futhark point of view, this grid is a two-dimensional array of integers that are either  $-1$  or  $1$ .

At any given discrete time step, the charge of a spin can be either positive or negative. A spin interacts only with its immediate neighbors, which makes Ising simulation a *stencil*. An atom prefers to have the same polarity as its neighbors, although it also has a small chance to flip polarity randomly, based on the temperature. This is the Monte Carlo aspect.

To update the grid, we compute for each spin  $c$  its corresponding *energy gradient*  $\Delta_e$ , as follows:

$$\Delta_e = 2c(u + d + l + r)$$

where  $u, d, l, r$  are the spins directly adjacent to  $c$  in the grid (we ignore diagonals). Further, for each spin we compute two random numbers,  $a$  and  $b$ , in the range  $(0, 1)$ . *It is very important that each spin receives its own  $a, b$ .* Then we compute the new value of  $c$  as

$$c' = \begin{cases} -c & \text{if } a < p \wedge (\Delta_e < -\Delta_e \vee b < e^{-\Delta_e \div t}) \\ c & \text{otherwise} \end{cases}$$

where  $0 \leq p \leq 1$  is the *sample rate*, and  $t \in \mathbb{R}$  is the *temperature*. Put in words,  $p$  is the fraction of spins that are candidates for flipping in a given time step. Of these, we flip those where it would locally reduce the energy of the system, or randomly, where the chance of random flips is proportional to the temperature. For more information on Ising models, I recommend the very readable 6-page article *The World in a Spin* by Brian Hayes<sup>3</sup>. However, we need not understand the physics to implement the model in Futhark.

For this exercise, you will be modifying the, `ising.fut` file in the code handout `ising-handout.tar.gz`, that contains the skeleton for an implementation of the 2D Ising model. Read the existing code carefully. Use `futhark pkg sync` to download the dependencies.

The code handout also comes with a program, `ising-tui.c`, that permits a console visualisation of the computation. Running the visualisation is not

---

<sup>3</sup><http://bit-player.org/wp-content/extras/bph-publications/AmSci-2000-09-Hayes-Ising.pdf>

required to solve the assignment; it's just a lot more satisfying. To run the visualisation, do

```
$ make ising-tui  
$ ./ising-tui
```

This may look more impressive (and take a longer time to converge) if you make your console font smaller.

#### Exercise 4.1: Generate initial state

The code handout defines the type of spins as follows:

```
type spin = i8
```

However, we also need to generate (potentially) distinct random numbers for every spin. For this exercise, we will use *one RNG state per spin*. Thus, the function for generating an initial grid state is:

```
entry random_grid (seed: i32) (h: i64) (w: i64)  
: ([h][w]rng_engine.rng, [h][w]spin) =  
...
```

Where `rng_engine.rng` is the type of RNG states. See the code comments for more information.

**Hint** Generate one-dimensional arrays of size  $n \times m$ , then use `unflatten` at the end.

#### Exercise 4.2: Computing $\Delta_e$

This is the stencil operation where we, for every spin, compute  $\Delta_e$  as a value of type `i8`:

```
entry deltas [h][w] (spins: [h][w]spin): [h][w]i8 =  
...
```

One question that must be answered for every stencil is how to handle the edges of the grid, where there are no neighbors. For the 2D Ising model, we pick the easy solution, and use *wraparound*, also known as a *torus world*. Simply put, when we go over one edge, we come out on the opposite edge. In Futhark, this is easily done with the `rotate` function. If `xss` is a two-dimensional array, then `rotate 1 xss` corresponds to rotating the array by one element vertically (along the first dimension), while `map (rotate (-1)) xss` rotates it by negative one element horizontally (along the second dimension).

**Hint** Use the `map2/map3/map4/map5` functions to map across multiple arrays simultaneously.

### Exercise 4.3: The step function

Define the step function, which computes one time step of the simulation:

```
entry step [h][w] (abs_temp: f32) (samplerate: f32)
    (rngs: [h][w]rng_engine.rng)
    (spins: [h][w]spin)
: ([h][w]rng_engine.rng, [h][w]spin) =
...
```

Note that it computes not just new spins, but also new RNG states.

### Exercise 4.4: Benchmarking

If you have defined the above functions correctly, then you should now be able to use the predefined `main` function, which creates a grid and runs a few steps of the simulation; returning the final grid at the end:

```
def main (abs_temp: f32) (samplerate: f32)
    (w: i64) (h: i64) (n: i32): [h][w]spin =
(loop (rngs, spins) = random_grid 1337 h w
  for _i < n do
    step abs_temp samplerate rngs spins).1
```

Show benchmarks that demonstrate how sequential versus parallel performance varies for different values of `w`, `h`, and `n`. Explain your results.

**Hint** Use `futhark bench` with the `--backend=c`, `--backend=multicore` and `--backend=opencl` (or `cuda`) options to switch code generators. Be aware that the Futhark compiler is not very good at generating sequential CPU code, and that stencils in particular are likely to have poor cache behaviour, so don't consider the CPU results indicative of the full power of your CPU.